

Continuum multiband exact diagonalization of twisted MoTe₂

An ED benchmark for CDW and fractional Chern-insulator neural states

Local numerical benchmark for the Physics-for-AI/QNS programme

3 September 2026
Version 3.0.0

Abstract

We establish a symmetry-resolved exact-diagonalization (ED) benchmark and a continuous-coordinate neural-Bloch experiment for twisted MoTe₂. The continuum calculation reproduces the one-five-band reference energies, while one-band 27-cell ED distinguishes the $\nu = 1/3$ charge-density-wave tower from the approximately three-fold Γ manifold at $\nu = 2/3$. A tracked three-state calculation gives Chern numbers near zero and two, but complete all-sector spectral isolation, large-gauge sewing and twist-grid convergence remain required before assigning a quantized Hall response.

The variational calculation follows the Luo-Fu Appendix-B neural-orbital construction in the full plane-wave continuum basis. Version 3 augments the exact- Γ no- M state by a trainable determinant mixer M_S , which linearly combines all 12 selected-momentum determinants obeying total-momentum conservation. Complete-wavefunction C_3 projection then resolves its three rotation sectors. Following the reference optimization convention, reported energies are final-ten means of the equilibrated persistent training chain; their RMS is a tail-stability diagnostic rather than an independent error bar. At $\nu = 1/3$, v3 no- M and the lowest outer sector P_0 give $E/N_e = -37.3902(287)$ and $-37.3458(516)$ meV, respectively. At $\nu = 2/3$, v3 no- M gives $-52.7841(394)$ meV and the lowest outer training sector is P_2 , with $E/N_e = -52.8211(477)$ meV, 0.0956 meV below five-band ED on the same training-chain comparison. Because this is a single-seed correlated optimization trace, it is not yet a controlled variational bound. A jointly batched complete-ratio estimator evaluates the five-band-projected 1-RDM for both v3 no- M states and all outer $m = 0, 1, 2$ sectors while retaining full graph, backflow and determinant recomputation for every particle replacement.

Part I

Exact-diagonalization benchmark

1 Motivation

Fractional Chern insulators (FCIs) in moiré semiconductors require simultaneous control of strong Coulomb correlations, band mixing and nonuniform Bloch-state quantum geometry. A quantum neural state (QNS) can reach larger systems and avoid a rigid few-band truncation, but a lower variational energy is not sufficient evidence of the correct phase. The reference calculation must also reproduce symmetry quantum numbers, the torus ground-state counting, momentum occupation and density correlations.

The present benchmark therefore separates two questions. First, the 3×3 multiband calculation tests the continuum Hamiltonian and the variational convergence from one to five retained bands, directly against Fig. 2 of Ref. [1]. Second, the one-band 27-site calculation repeats the phase diagnostics of Supplementary Fig. 5: a Γ, K, K' CDW tower at $\nu = 1/3$ versus an approximately three-fold Γ manifold at $\nu = 2/3$ [1]. The latter momentum counting is supplemented below by flux spectral flow and a gauge-invariant many-body Chern number.

2 Continuum model and quantum geometry

For a fixed spin/valley sector, the single-hole Hamiltonian is

$$H_0 = \begin{pmatrix} \frac{\hbar^2(-i\nabla - \kappa_+)^2}{2m^*} + V_1(\mathbf{r}) & t(\mathbf{r}) \\ t^*(\mathbf{r}) & \frac{\hbar^2(-i\nabla - \kappa_-)^2}{2m^*} + V_2(\mathbf{r}) \end{pmatrix}. \quad (1)$$

The lowest moiré harmonics and plane-wave convention follow Refs. [5, 1]. We use

$$(V, w, \phi, m^*) = (11.2 \text{ meV}, 13.3 \text{ meV}, -91^\circ, 0.62m_e), \quad (2)$$

monolayer lattice constant $a_0 = 0.352$ nm, and eight radial reciprocal shells. The resulting Bloch Hamiltonian

has 61 plane waves per layer. Figure 1 compares the first ten moiré bands at $\theta = 2.6^\circ$ and 3.0° .

For a nondegenerate band n , the quantum geometric tensor is evaluated from velocity matrix elements,

$$Q_{\alpha\beta}^{(n)}(\mathbf{k}) = \sum_{m \neq n} \frac{\langle n | v_\alpha | m \rangle \langle m | v_\beta | n \rangle}{[\varepsilon_n(\mathbf{k}) - \varepsilon_m(\mathbf{k})]^2}, \quad (3)$$

with $g_{\alpha\beta} = \text{Re } Q_{\alpha\beta}$ and $\Omega = -2 \text{Im } Q_{xy}$. Bands 1–3 are calculated on a 49-point linear grid in each reciprocal direction, restricted to the hexagonal Brillouin zone, and smoothly triangulated only for visualization. The shared colorbars retain physical one-dimensional numerical ticks. Values outside the 1.5–98.5 percentile interval are clipped in the rendering, not in the saved data.

3 Projected interaction and ED symmetries

The periodically continued Coulomb interaction is projected into N_b Bloch bands,

$$H = \sum_{n\mathbf{k}} \varepsilon_{n\mathbf{k}} c_{n\mathbf{k}}^\dagger c_{n\mathbf{k}} + \sum_{ab,cd} V_{cd,ab} c_{c\mathbf{k}}^\dagger c_{d\mathbf{k}}^\dagger c_{b\mathbf{k}} c_{a\mathbf{k}}, \quad (4)$$

where $V(\mathbf{q}) = 2\pi e^2/(\epsilon|\mathbf{q}|)$ and the uniform $\mathbf{q} = 0$ component is removed by the neutralizing background. Eight radial reciprocal shells are retained in the form-factor sum.

Charge- $U(1)$ is enforced by fixed hole number, while spin- $U(1)$ is enforced by restricting to the fully polarized sector. Translation symmetry first decomposes the Hilbert space into many-body momentum blocks $\mathcal{H}_K$. At generic momentum, a rotation maps $\mathcal{H}_K$ to the distinct block $\mathcal{H}_{C_n K}$, so point-group symmetry is used only to copy spectra within a momentum orbit. A further C_3/C_6 character decomposition is available only when $C_n K = K + G$, i.e. when the rotation belongs to the little group of K .

The phase-diagnostic cluster has real-space supercell matrix

$$S_{27} = \begin{pmatrix} 3 & -3 \\ 3 & 6 \end{pmatrix}, \quad \det S_{27} = 27, \quad (5)$$

and is therefore the $3\sqrt{3} \times 3\sqrt{3}$ C6 torus rather than a rectangular 3×9 mesh. Its momentum group is written as $\mathbb{Z}_3 \times \mathbb{Z}_9$, with plot index

$$I_k = k_1 + N_1 k_2, \quad N_1 = 3, \quad I_k = 0, \dots, 26. \quad (6)$$

In this convention Γ, K', K occur at $I_k = 0, 9, 18$. The one-band Hilbert space contains $\binom{27}{9} = 4,686,825$ states at $\nu = 1/3$. At $\nu = 2/3$, an exactly verified particle–hole transformation maps 18 particles to 9 holes without changing the spectrum.

4 Multiband benchmark on the 3×3 cluster

Figure 2 summarizes the 3×3 calculation at $\theta = 3^\circ$, $\epsilon = 5$. The ground-state energy decreases monotonically as the retained band space grows. The values in Table 1 reproduce the available $\nu = 2/3$ points in Fig. 2 of Ref. [1] within 0.036 meV per hole. Remote-band mixing is visibly stronger at $\nu = 2/3$.

Table 1: Ground-state energy per hole in meV at $\theta = 3^\circ$, $\epsilon = 5$ on the 3×3 cluster.

N_b	$\nu = 1/3$	$\nu = 2/3$	paper Fig. 2
1	−36.3296	−47.1561	−47.15
2	−37.1158	−50.0774	−50.06
3	−37.3518	−51.8605	−51.83
4	−37.3878	−52.5379	—
5	−37.3932	−52.7255	−52.69

The middle column of Fig. 2 shows the lowest energy in every momentum sector for one-, three- and five-band ED. Energies are plotted as $E_K - E_{\text{GS}}$ in meV, without division by particle number. The right column reports the additive band-resolved weight

$$w_b = \frac{\langle N_b \rangle}{N_h}, \quad \sum_{b=1}^{N_b} w_b = 1. \quad (7)$$

For five bands, the weights are

$$\nu = 1/3 : (0.96469, 0.02953, 0.00523, 0.00049, 0.00005), \quad (8)$$

$$\nu = 2/3 : (0.88952, 0.07591, 0.02535, 0.00763, 0.00159). \quad (9)$$

These are normalized band populations, not overlaps between separately optimized truncated Hilbert spaces.

5 Reproduction of Supplementary Fig. 5

Figure 3 repeats the one-band 27-site benchmark at $\theta = 2.6^\circ$, $\epsilon = 5$, and adds the ground-manifold averaged $n(\mathbf{k})$. At $\nu = 1/3$, the three global low states occur at Γ, K, K' with

$$(E_\Gamma, E_K, E_{K'}) - E_{\text{GS}} = (0, 0.573, 0.573) \text{ meV}. \quad (10)$$

The next state is 1.476 meV above the top of this manifold. The K, K' states fold to Γ in the Brillouin zone of the tripled CDW unit cell, matching the interpretation of Ref. [1].

At $\nu = 2/3$, the three lowest states all lie in the Γ sector,

$$(E_{\Gamma,0}, E_{\Gamma,1}, E_{\Gamma,2}) - E_{\text{GS}} = (0, 0.108, 0.298) \text{ meV}. \quad (11)$$




Figure 1: Continuum bands and quantum geometry. (a,b) The first ten moiré bands along Γ - K - M - Γ at $\theta = 2.6^\circ$ and 3.0° ; bands 1–3 are colored blue, orange and green, respectively. (c–e) Berry curvature and (f–h) trace quantum metric for bands 1–3 at $\theta = 2.6^\circ$ across the hexagonal Brillouin zone. Both rows use a common physical color scale with explicit numeric ticks.

Thus the finite-size threefold splitting is small but explicitly nonzero. The gap between the highest manifold state and the next global level is 3.482 meV. This reproduces the torus counting required for the $\nu = 2/3$ FCI/FQAH state.

The absolute energy origins used to construct the two spectra in Fig. 3 are listed explicitly in Table 2. They are the global minima over all 27 translation sectors; the plotted ordinate is obtained by subtracting the corresponding total energy in the fourth column, without dividing the excitation energy by particle number.

Table 2: Absolute one-band, 27-cell ground-state energies used as the zero of the Fig. 3 spectra.

ν	N_h	three-state manifold	E_{GS} (meV)	E_{GS}/N_h (meV)
1/3	9	Γ, K, K'	-428.359163	-47.595463
2/3	18	$\Gamma_0, \Gamma_1, \Gamma_2$	-1034.203144	-57.455730

For each filling, $n(\mathbf{k})$ and the projected static structure factor are averaged over the three-state ground manifold,

$$n(\mathbf{k}) = \langle c_{\mathbf{k}}^\dagger c_{\mathbf{k}} \rangle, \quad (12)$$

$$S(\mathbf{q}) = \frac{1}{N_h} \langle \rho(-\mathbf{q}) \rho(\mathbf{q}) \rangle, \quad \mathbf{q} \neq 0. \quad (13)$$

The connected value $S(0)$ is set to zero. The $n(\mathbf{k})$ sums are exactly 9 and 18. For Fig. 3, the discrete $n(\mathbf{k})$ markers are

placed over an interpolated, weak band-1 Berry-curvature background. The $S(\mathbf{q})$ markers in both rows are placed over the same interpolated band-1 trace-condition deviation $\text{Tr } g_1(\mathbf{q}) - |\Omega_1(\mathbf{q})|$. The background palettes match Fig. 1, while the observable markers use brighter independent color scales.

At $\nu = 1/3$, $S(K) = S(K') = 1.057$, producing the Bragg peaks characteristic of CDW order. At $\nu = 2/3$, the maximum at the same momenta is only 0.197 and is comparable to the surrounding finite- q response; no CDW Bragg peak is present. Together with the three Γ states, this is consistent with the FQAH liquid signature reported in Supplementary Fig. 5.

6 Twisted-boundary spectra and many-body Berry curvature

Let $\mathbf{a}_1, \mathbf{a}_2$ be the primitive moiré real-space vectors and let $\mathbf{b}_a \cdot \mathbf{a}_b = 2\pi\delta_{ab}$. The C_6 -compatible 27-cell ED torus is generated by the columns of

$$S = \begin{pmatrix} 3 & -3 \\ 3 & 6 \end{pmatrix}, \quad |\det S| = 27, \quad (14)$$

$$\mathbf{L}_1 = 3\mathbf{a}_1 + 3\mathbf{a}_2,$$

$$\mathbf{L}_2 = -3\mathbf{a}_1 + 6\mathbf{a}_2.$$




Figure 2: Multiband ED on the 3×3 cluster. Top row: $\nu = 1/3$; bottom row: $\nu = 2/3$. (a,d) Ground-state energy per hole versus retained bands; filling labels are placed in the upper-right blank region. (b,e) The lowest three levels in every momentum sector for one-, three- and five-band ED, plotted as $E - E_{\text{GS}}$ in meV. Larger opaque markers denote sector ground states; the next two levels are smaller and semitransparent. (c,f) Normalized band-resolved ground-state weights. Charge- $U(1)$, spin- $U(1)$ and translations are exact; C_3/C_6 equivalence accelerates the largest five-band sectors.

For every particle i , the many-body twisted boundary condition is

$$\Psi(\dots, \mathbf{r}_i + \mathbf{L}_a, \dots) = e^{i\phi_a} \Psi(\dots, \mathbf{r}_i, \dots), \quad \phi_a \in [0, 2\pi]. \quad (15)$$

Writing a one-particle momentum as $\mathbf{k} = B\mathbf{f}$, with $B = (b_1, b_2)$, Eq. (15) gives $S^T \mathbf{f} = \mathbf{m} + \phi/(2\pi)$. Hence the 27 allowed momenta are

$$\mathbf{f}_{\mathbf{m}}(\phi) = S^{-T} \left(\mathbf{m} + \frac{\phi}{2\pi} \right), \quad \mathbf{m} \in \mathbb{Z}^2 / S^T \mathbb{Z}^2, \quad (16)$$

$$\mathbf{k}_{\mathbf{m}}(\phi) = \mathbf{k}_{\mathbf{m}}(0) + \delta \mathbf{k}, \quad \delta \mathbf{k} = BS^{-T} \frac{\phi}{2\pi}. \quad (17)$$

The code enumerates this quotient group with exact integer residues, so precisely 27 inequivalent momenta are retained and every two-body momentum addition is performed modulo $S^T \mathbb{Z}^2$. The torus is therefore implemented periodically without rounding Cartesian momenta.

At every (ϕ_1, ϕ_2) , the continuum Hamiltonian is solved anew at all $\mathbf{k}_{\mathbf{m}}(\phi)$, after which the one-band Coulomb form factors and projected interaction are rebuilt from the shifted Bloch spinors. A projected scattering term is retained only when $\mathbf{k}_1 + \mathbf{k}_2 = \mathbf{k}_3 + \mathbf{k}_4$ in the finite translation group (equivalently up to moiré and supercell reciprocal vectors). Charge- $U(1)$ and translations remain

exact. A generic twist breaks C_3/C_6 , so rotational orbit copying is not used away from little-group-invariant twists. At $\nu = 2/3$, the 18-electron Hamiltonian is evaluated exactly through its nine-hole particle-hole transform, including its constant and induced one-body terms; this leaves the boundary condition and spectrum unchanged.

Because the three lowest states are quasi-degenerate, an individual-eigenvector Berry phase is gauge dependent. We instead use the determinant link of the full three-state subspace,

$$M_{\mu}^{ab}(\phi) = \langle \Psi_a(\phi) | \Psi_b(\phi + \Delta_{\mu}) \rangle, \quad (18)$$

$$U_{\mu}(\phi) = \frac{\det M_{\mu}(\phi)}{|\det M_{\mu}(\phi)|}, \quad (19)$$

$$\Phi_{12}(\phi) = \arg[U_1(\phi)U_2(\phi + \Delta_1)U_1^*(\phi + \Delta_2)U_2^*(\phi)]. \quad (20)$$

The overlap in M_{μ} is not a bare dot product of ED coefficient vectors, because the one-body Bloch basis itself changes with twist. For an electron configuration I , its coefficient overlap is weighted by

$$w_I = \prod_{p \in I} \langle u_p(\phi) | u_p(\phi + \Delta_{\mu}) \rangle. \quad (21)$$

For the nine-hole representation, the code uses the exactly




Figure 3: 27-site one-band reproduction of Supplementary Fig. 5. Top row: $\nu = 1/3$; bottom row: $\nu = 2/3$. (a,d) Momentum-resolved spectra, plotted as $E - E_{\text{GS}}$ in meV. Red markers denote the CDW or FCI three-state manifold. (b,e) Ground-manifold averaged $n(\mathbf{k})$ over weak, interpolated band-1 Berry curvature. (c,f) Projected $S(\mathbf{q})$ over the identical weak, interpolated band-1 $\text{Tr } g_1 - |\Omega_1|$ background. Bright marker colorbars encode the observables independently. The sharp zone-corner peaks at $\nu = 1/3$ and their absence at $\nu = 2/3$ distinguish the CDW and FCI responses.

equivalent electron complement,

$$w_H = \frac{\prod_{p=1}^{27} \langle u_p(\phi) | u_p(\phi + \Delta_\mu) \rangle}{\prod_{h \in H} \langle u_h(\phi) | u_h(\phi + \Delta_\mu) \rangle}. \quad (22)$$

The non-Abelian FHS curvature and discrete Chern number are

$$\begin{aligned} \mathcal{F}_{\phi_1 \phi_2} &= \frac{\Phi_{12}}{(\Delta\phi)^2}, \\ C_{\text{bundle}} &= \frac{1}{2\pi} \sum_{i,j} \Phi_{12}(\phi_{ij}), \\ \Delta\phi &= \frac{2\pi}{6}. \end{aligned} \quad (23)$$

We solve and cache the 7×7 grid endpoints as well as the 6×6 plaquettes. Links crossing a 2π boundary are consequently evaluated from endpoint Bloch overlaps, rather than by identifying raw ED eigenvectors. A separate analytic large-gauge sewing-matrix test is nevertheless still needed for a formal convergence claim. This single grid is a discrete diagnostic, not a twist-grid convergence study.

Figure 4 shows the two axial flux spectra and the curvature map. For $\nu = 1/3$, the bundle contains the lowest states from Γ, K, K' . Its Chern number is

$$C_{\text{bundle}}^{1/3} = 9.3 \times 10^{-7} \simeq 0, \quad (24)$$

and the smallest sampled separation between the three-state manifold and the nearest state in these three sectors is 1.476 meV. For $\nu = 2/3$, the three states all belong to Γ and

$$C_{\text{tracked}}^{2/3} = 1.99999975 \simeq 2. \quad (25)$$

If this tracked subspace is subsequently shown to be the globally isolated ground-state bundle, the conditional Hall response is $\sigma_{xy} = C_{\text{tracked}} e^2 / (3h) = 2e^2 / (3h)$. The minimum same- Γ separation is 5.144 meV, the minimum link singular value is 0.958 at $\nu = 2/3$, and the minimum tracked-sector overlap is 0.964 at $\nu = 1/3$. These checks show that the sampled links are nonsingular, but they do not prove isolation from the complete spectrum. Missing checks are the lowest relevant levels in all 27 translation sectors at every twist, the global $\min_\phi (E_4 - E_3)$, an explicit analytic 2π large-gauge sewing test, and $6 \times 6 \rightarrow 9 \times 9 \rightarrow 12 \times 12$ grid convergence.

7 ED benchmark protocol for QNS

The validated workflow is:

1. solve the plane-wave continuum Hamiltonian and verify shell convergence;

27-cell one-band ED under twisted boundaries ($\theta = 2.6^\circ$, $\varepsilon = 5$)




Figure 4: Twisted-boundary ED and many-body topology on the 27-cell one-band cluster. Top row: $\nu = 1/3$; bottom row: $\nu = 2/3$. (a,d) Spectral flow versus ϕ_1 at $\phi_2 = 0$. (b,e) Spectral flow versus ϕ_2 at $\phi_1 = 0$. At every twist, the lowest three energies are red scatter points. Gray points show five additional $\nu = 2/3$ levels and nine additional $\nu = 1/3$ levels sampled from the three CDW sectors. (c,f) Non-Abelian curvature $\mathcal{F}_{\phi_1\phi_2}$ on the original 6×6 primitive twist grid; white dots mark the FHS plaquette centers. The raw plaquette-flux sums give Chern number zero for the CDW and two for the tracked FCI manifold.

2. calculate band dispersion, Berry curvature and quantum metric on the same Bloch convention;
3. project the neutral-background Coulomb interaction with reciprocal-vector form factors;
4. fix charge- $U(1)$ and spin- $U(1)$, then block by translation; use point-group orbit equivalence at generic K , and split by C_3/C_6 character only when the rotation belongs to the little group of K ;
5. use sparse Hermitian Lanczos independently in each required sector;
6. compare energies, sector counting, normalized band weights, $n(\mathbf{k})$ and $S(\mathbf{q})$;
7. insert boundary flux, track the isolated ground manifold and integrate its non-Abelian Berry curvature;
8. only then benchmark the QNS energy, variance and state overlaps against ED.

Numerical regression checks support every stage. The general C_6 -supercell interaction agrees with the earlier rectangular 3×3 implementation to relative error 3.3×10^{-16} . Direct electron and particle-hole spectra agree

sector by sector to 3.4×10^{-13} meV. The finite plane-wave cutoff produces a maximum one-body C_6 mismatch of 8.0×10^{-6} meV, and the pre-symmetrization interaction Hermiticity error is 7.8×10^{-9} . The current complete 51-test regression suite passes; the topology outputs additionally record link conditioning, Lanczos residuals and interaction Hermiticity at every twist.

Part II

Neural many-body Bloch wavefunctions

8 Motivation and scope

The second calculation tests the neural many-body Bloch construction of Ref. [1] directly in the continuum model benchmarked in Part I. The purpose is not to fit coefficients in the one-band ED Hilbert space. Instead, a graph network dresses continuum Bloch orbitals in coordinate space and the resulting generalized Slater determinant is optimized against the complete interacting

continuum Hamiltonian. This can generate remote-band content even though the fixed one-body input is the mean-field first band.

This revision follows Appendix B, Eqs. (16)–(23), of Ref. [1]. In particular, the orbital transformation is the direct complex output

$$J_{\mathbf{k}i}(X) = \text{MLP}([\mathbf{k}, V_i^t]) \in \mathbb{C}, \quad Q_{\mathbf{k}i} = J_{\mathbf{k}i} \varphi_{\mathbf{k}i}. \quad (26)$$

There is no constant-one skip, residual parameterization, or factor 0.2; the implemented transformation is the direct $J_{\mathbf{k}i}$ above. The current source additionally implements three audit corrections: full- M uses an unbiased complex random initialization; the node aggregation is the literal $\sum_j m_{ij}$ in Eq. (19), including $j = i$; and the random auxiliary features h_i^0, h_{ij}^0 are explicitly retained in the trainable parameter tree. Version 3 further replaces the v2 single selected-momentum no- M determinant by a trainable M_S matrix that linearly combines every selected-momentum Slater determinant in the exact total- Γ sector. Earlier one-hot- M , no-self-message and v2 single-minor checkpoints are excluded from the v3 comparison.

Both v3 no- M runs and all six v3 outer- C_3 branches are complete. The internal-equivariance construction investigated in v2 is omitted from the present results because its threefold evaluation was expensive and did not improve the controlled energy comparison. Quoted training energies use the final ten persistent-chain measurements; the corresponding RMS is a stability diagnostic rather than an independent error bar.

9 Configuration, fixed Bloch basis, and conventions

A walker is a first-quantized configuration

$$X = \{(\mathbf{r}_i, l_i)\}_{i=1}^{N_e}, \quad \mathbf{r}_i \in [0, 1]^2, \quad l_i^{\text{code}} \in \{0, 1\}. \quad (27)$$

The fractional coordinate is converted to a Cartesian coordinate on the 3×3 torus. The layer input used by the network is

$$s_i = 2l_i^{\text{code}} - 1 \in \{-1, +1\}, \quad (28)$$

which is the $l_i = \pm 1$ convention of Ref. [1]. Particle labels are only row labels of the determinant and graph; no ordered occupation string is used.

The nine mesh momenta and plane-wave orbitals are

$$\mathbf{k}_{m_1 m_2} = \frac{m_1}{3} \mathbf{b}_1 + \frac{m_2}{3} \mathbf{b}_2, \quad m_1, m_2 = 0, 1, 2, \quad (29)$$

$$\phi_{\mathbf{k}}^0(\mathbf{r}, l) = \sum_{\mathbf{G}} u_{l\mathbf{k}\mathbf{G}}^l e^{i(\mathbf{k}+\mathbf{G})\cdot\mathbf{r}}. \quad (30)$$

The reciprocal set contains the same 61 $\mathbf{G}$ vectors per layer used for the single-particle bands in Fig. 1. The coefficients $u_{l\mathbf{k}\mathbf{G}}^l$ are obtained once by diagonalizing the continuum mean-field Hamiltonian. In JAX they are constants, and in PyTorch they are registered buffers. They are absent from the trainable parameter tree. Thus every

run uses band 1 as its mean-field orbital source; neither ED amplitudes nor trainable plane-wave coefficients are used.

10 Appendix-B graph network

10.1 Initial physical and auxiliary features

The three reciprocal harmonics used at each vertex are $\mathbf{g}_a \in \{\mathbf{b}_1, \mathbf{b}_2, \mathbf{b}_1 + \mathbf{b}_2\}$. The immutable node feature is

$$v_i^0 = [\sin(\mathbf{g}_a \cdot \mathbf{r}_i), \cos(\mathbf{g}_a \cdot \mathbf{r}_i), s_i] \in \mathbb{R}^7. \quad (31)$$

For every ordered pair (i, j) , the immutable edge feature is

$$e_{ij}^0 = [\sin(\mathbf{k} \cdot (\mathbf{r}_i - \mathbf{r}_j)), \cos(\mathbf{k} \cdot (\mathbf{r}_i - \mathbf{r}_j)), s_i] \in \mathbb{R}^{19}, \quad (32)$$

where all nine mesh momenta are included. In particular, the edge layer channel is the source layer s_i , exactly as in Eq. (16) of Ref. [1]; no additional layer channels are appended.

The complete initial features are

$$V_i^0 = (v_i^0, h_i^0) \in \mathbb{R}^{32}, \quad E_{ij}^0 = (e_{ij}^0, h_{ij}^0) \in \mathbb{R}^{32}. \quad (33)$$

The code therefore uses 25 auxiliary node components and 13 auxiliary edge components. Appendix B explicitly calls h_i^0 and h_{ij}^0 “randomly initialized parameters.” They are therefore leaves of the trainable JAX pytree (and `nn.Parameter` objects in PyTorch), so natural-gradient updates act on them together with the other network parameters. The present code uses one random vector shared over particle labels and one shared over ordered pairs; sharing preserves permutation equivariance. Each component is initialized from $\mathcal{N}(0, 1/32)$. The paper does not specify this distribution or sharing rule, so both remain explicit local conventions. There is no learned input embedding before Eq. (36).

10.2 Two message-passing iterations

For iteration $a = 1, 2$,

$$Q_{ij}^a = W_Q^a E_{ij}^{a-1}, \quad (34)$$

$$K_{ij}^a = W_K^a E_{ij}^{a-1}, \quad (35)$$

$$m_{ij}^a = F^a \left(\sum_n Q_{in}^a \otimes K_{nj}^a \right) \odot G^a(E_{ij}^{a-1}), \quad (36)$$

$$h_i^a = h_i^{a-1} + U^a \left(\left[\sum_j m_{ij}^a, V_i^{a-1} \right] \right), \quad (37)$$

$$h_{ij}^a = h_{ij}^{a-1} + H^a([m_{ij}^a, E_{ij}^{a-1}]), \quad (38)$$

$$V_i^a = (v_i^0, h_i^a), \quad (39)$$

$$E_{ij}^a = (e_{ij}^0, h_{ij}^a). \quad (40)$$

The outer product in Eq. (36) is flattened before F^a . F^a, G^a, H^a, U^a are one-layer real affine maps followed by

exact GELU. The two iterations have independent parameters. No division by N_e , no neighbor-count normalization, and no LayerNorm is applied. Self edges remain available to both the edge update and node aggregation. This is the literal $\sum_j$ printed in Luo–Fu Eq. (19), rather than the “all other particles” neighborhood convention of Ref. [3]. The literal physical features v_i^0, e_{ij}^0 are reinserted after every iteration.

11 Backflow, orbital transformation, and determinant

The complex backflow matrix is implemented as a real, bias-free $32 \rightarrow 4$ linear map. Its first and last two outputs are the real and imaginary parts of a two-component displacement,

$$\delta \mathbf{r}_i(X) = W V_i^2, \quad \tilde{\mathbf{r}}_i = \mathbf{r}_i + \delta \mathbf{r}_i. \quad (41)$$

There is no tanh, clipping, or external backflow scale. The fixed mean-field coefficients are evaluated only at this transformed coordinate,

$$\varphi_{\mathbf{k}i}(X) = \sum_{\mathbf{G}} u_{1\mathbf{k}\mathbf{G}}^i \exp[i(\mathbf{k} + \mathbf{G}) \cdot (\mathbf{r}_i + \delta \mathbf{r}_i(X))]. \quad (42)$$

Consequently, backflow does not update $u_{1\mathbf{k}\mathbf{G}}^i$. The factors $\exp[i(\mathbf{k} + \mathbf{G}) \cdot \delta \mathbf{r}_i]$ are generated only by changing the coordinate argument.

The neural orbital transformation uses the Cartesian two-component mesh momentum and the final node feature,

$$[V_i^2, k_x, k_y] \in \mathbb{R}^{34} \xrightarrow{34 \rightarrow 32} \text{GELU} \xrightarrow{32 \rightarrow 2} (J_{\mathbf{k}i}^R, J_{\mathbf{k}i}^I), \quad (43)$$

$$J_{\mathbf{k}i} = J_{\mathbf{k}i}^R + i J_{\mathbf{k}i}^I, \quad Q_{\mathbf{k}i} = J_{\mathbf{k}i} \varphi_{\mathbf{k}i}. \quad (44)$$

The direct complex output is not converted to $1 + \alpha \Delta J$. Equation (44) is implemented once: J is not included inside φ and then multiplied again. A unit test that sets $J_{\mathbf{k}i} = 1/2$ verifies $\det Q = 2^{-N_e} \det \varphi$ to numerical precision.

For the Luo–Fu full- M state,

$$\Phi_{\gamma i}(X) = \sum_{\mathbf{k}} M_{\gamma \mathbf{k}} Q_{\mathbf{k}i}(X), \quad (45)$$

$$\Psi_M(X) = \det_{\gamma i} \Phi_{\gamma i}(X), \quad M \in \mathbb{C}^{N_e \times 9}. \quad (46)$$

All graph, W , J , and complex M parameters are trainable. Only one final generalized Slater determinant is used. Because M is trainable and dense, the Cauchy–Binet expansion of $\det(MQ)$ generally contains several total-momentum sectors.

Version 3 removes the dense single-particle matrix M while retaining a trainable many-body mixing within the exact total-momentum sector. Let $\mathcal{C}_\Gamma^{(N_e)}$ be the set of canonical N_e -element subsets I of the nine mesh momenta

satisfying $\sum_{\mathbf{k} \in I} \mathbf{k} = \Gamma$ modulo a primitive reciprocal vector. The v3 no- M state is

$$\Psi_\Gamma^{\text{v3}}(X) = \sum_{I \in \mathcal{C}_\Gamma^{(N_e)}} (M_S)_{1I} \det_{\mathbf{k} \in I, i} Q_{\mathbf{k}i}(X), \quad M_S \in \mathbb{C}^{N_D \times |\mathcal{C}_\Gamma^{(N_e)}|}. \quad (47)$$

Both $N_e = 3$ and $N_e = 6$ have 12 such subsets on the 3×3 momentum mesh. The determinant-mixing matrix M_S is initialized with a split-real LeCun-normal convention and optimized by the same stochastic-reconfiguration update as the graph, backflow and orbital-transformation parameters. Here $N_D = 1$, so M_S has shape 1×12 and contributes 24 real trainable parameters. Thus v3 fixes total momentum exactly to Γ but does not freeze the many-body momentum content to one bare Slater determinant. The v2 state, retained only as a historical baseline, kept the single subset that minimized the sum of bare band-1 energies: $\{(2, 0), (2, 1), (2, 2)\}/3$ for $N_e = 3$ and $\{(1, 0), (1, 1), (1, 2), (2, 0), (2, 1), (2, 2)\}/3$ for $N_e = 6$.

The v3 matrix M_S is not the Luo–Fu dense matrix M . A dense M mixes momenta at the one-particle orbital level and its Cauchy–Binet expansion generally spans several total-momentum sectors. In contrast, Eq. (47) uses M_S only at the many-body determinant level and contains exactly the 12 minors in the Γ sector. Each minor still uses the fully configuration-dependent $Q_{\mathbf{k}i}(X)$, so all graph, backflow and orbital transformations are recomputed before the linear combination is formed.

12 Complete-wavefunction outer C_3 projection

The outer construction applies no internal Reynolds average. In v3 its raw state is the periodic, exact-total- Γ sector mixture

$$D_\theta^{\text{v3}}(X) = \sum_{I \in \mathcal{C}_\Gamma^{(N_e)}} (M_S)_{1I} \det_{\mathbf{k} \in I, i} Q_{\mathbf{k}i}(X), \quad (48)$$

with an unrestricted Luo–Fu message-passing backflow and J network. Thus neither an individual displacement nor an individual generalized orbital is forced to be a C_3 eigenobject, while the complete raw state remains in the translation sector $\mathbf{K}_{\text{tot}} = \Gamma$.

The physical continuum action includes the layer-dependent Bloch-gauge sewing. With the code’s active-rotation convention, define

$$(C^a D_\theta^{\text{v3}})(X) = G_a(X) D_\theta^{\text{v3}}(C_3^a X), \quad (49)$$

where $G_0 = 1$ and G_a is the accumulated product of the physical layer-gauge factor at each preceding point on the configuration orbit. For $\omega = \exp(2\pi i/3)$, the complete-

state projector is

$$\begin{aligned} \Psi_{m,\theta}^{v3}(X) &= P_m D_\theta^{v3}(X) \\ &= \frac{1}{3} \sum_{a=0}^2 \omega^{-ma} G_a(X) \\ &\quad \times D_\theta^{v3}(C_3^a X), \quad m = 0, 1, 2. \end{aligned} \quad (50)$$

The three independently optimized sectors have the same 78,124 real parameters as v3 no M ; projection adds no trainable parameter. Every raw branch is a complete graph/determinant evaluation with shared parameters.

Crucially, all stochastic quantities use Eq. (50), not one raw branch. Metropolis ratios are $|\Psi_m(X')/\Psi_m(X)|^2$, every kinetic or tunnelling local-energy term contains the complete ratio $\Psi_m(X')/\Psi_m(X)$, and SR uses $\mathcal{O}_\mu(X) = \partial_{\theta_\mu} \log \Psi_m(X)$. Setting `outer_c3_projector=False` recovers the no- M wavefunction exactly.

Before production, 64 random configurations gave relative projection norms 0.5411, 0.5929 and 0.6410 for $m = 0, 1, 2$, so no sector was annihilated. The maximum character errors in $\Psi_m(C_3 X)/\Psi_m(X) = \omega^m$ were 6.13×10^{-14} , 6.39×10^{-14} and 5.02×10^{-14} . All idempotency/orthogonality residuals for $P_m P_n = \delta_{mn} P_m$ were below 7.55×10^{-15} , and disabling the projector changed the no- M output by exactly zero. These tests isolate the new outer-projection logic from the

13 Initialization and trainable parameter count

Every affine neural-network kernel uses JAX LeCunNormal, `std = 1/sqrt(fanin)`, with zero biases. The auxiliary h^0 convention was specified above. For full- M , the real and imaginary components are independently initialized with half-scale LeCunNormal,

$$\mathbb{E}[(\Re M_{\gamma k})^2] = \mathbb{E}[(\Im M_{\gamma k})^2] = \frac{1}{2N_k}, \quad \mathbb{E}|M_{\gamma k}|^2 = \frac{1}{N_k}. \quad (51)$$

Thus every momentum minor is accessible at step zero and there is no preferred Γ subset. Ref. [1] states global LeCunNormal network initialization but does not separately specify M ; the complex extension and variance split above are therefore the present user-requested convention, not an attributed statement from the paper.

The mean-field orbital source is always band 1. However, because W and J themselves are LeCunNormal rather than forced to $W = 0, J = 1$, the step-zero many-body state is not claimed to equal the undressed bare Slater determinant pointwise. This distinction is recorded to make the initialization auditable.

Table 3: Trainable real parameter counts. The v3 no- M and outer states use the 1×12 complex determinant mixer M_S .

filling	N_e	full M	no $M+M_S, \Gamma$	outer P_m+M_S
1/3	3	78,154	78,124	78,124
2/3	6	78,208	78,124	78,124

Relative to v2 no- M , $M_S \in \mathbb{C}^{1 \times 12}$ adds 24 real parameters. The external P_m projector shares the same raw parameters among all three rotated branches and adds no trainable parameter.

14 Hamiltonian and local-energy estimator

The VMC objective uses the complete continuum Hamiltonian at $\theta = 3^\circ$, screening parameter $\epsilon = 5$, and a 3×3 torus. It is not the one-band or five-band projected ED Hamiltonian. The local energy contains the kinetic term, layer-dependent moiré potential, spatially varying inter-layer tunnelling, and periodic Coulomb interaction,

$$E_{\text{loc}}(X) = \frac{(H\Psi)(X)}{\Psi(X)} = E_{\text{kin}} + E_{\text{moiré}} + E_{\text{tun}} + E_C. \quad (52)$$

The Coulomb sum uses reciprocal-shell cutoff 8, dielectric $\epsilon = 5$, and removes the $\mathbf{q} = 0$ neutral-background component. Tunnelling is evaluated by flipping one layer index and multiplying the exact wavefunction ratio.

Cartesian derivatives of the complete complex state use the three-point central stencil

$$\partial_a \Psi / \Psi \simeq \frac{\Psi(X + h\hat{a}) - \Psi(X - h\hat{a})}{2h\Psi(X)}, \quad (53)$$

$$\partial_a^2 \Psi / \Psi \simeq \frac{\Psi(X + h\hat{a}) + \Psi(X - h\hat{a}) - 2\Psi(X)}{h^2\Psi(X)}. \quad (54)$$

Here $h = 1.5 \times 10^{-3} a_M$. Periodic wrapping is applied after every coordinate displacement. A previous $h, h/2, h/4$ audit is not reused as a validation of the new checkpoints; it must be repeated after convergence.

14.1 Raw-coordinate and full-wavefunction Hamiltonian audit

The production JAX path is `src/jax_neural_block.py:184-248`; the independent PyTorch reference is `src/continuum_vmc.py:96-206`. Both implementations first form $\mathbf{r}_i$ from the input *raw positions*. Consequently,

$$V_M(X) = -2V \sum_{i,a} \cos(\mathbf{g}_a \cdot \mathbf{r}_i + \phi_{li}), \quad (55)$$

$$V_C(X) = \sum_{i < j} \sum_{\mathbf{q} \neq 0} v_{\mathbf{q}} \cos[\mathbf{q} \cdot (\mathbf{r}_i - \mathbf{r}_j)] \quad (56)$$

never see $\delta \mathbf{r}_i$. For tunnelling, $t(\mathbf{r}_i)$ is evaluated at that same raw coordinate, while its numerator is the complete

$\Psi(X_{l_i \rightarrow 1-l_i})$ obtained by a new `logpsi` call. The kinetic ratios similarly use new complete calls,

$$R_{i\mu}^{\pm} = \frac{\Psi(X_{\mathbf{r}_i \rightarrow \mathbf{r}_i \pm h\hat{\mu}})}{\Psi(X)}, \quad (57)$$

so each displaced call rebuilds the graph and recomputes every particle's $\delta\mathbf{r}_j$, $J_{\mathbf{k}j}$, the MQ contraction, and the determinant. No backflow displacement is cached by the Hamiltonian.

As a final energy-zero and convention check, a fixed-state regression sets $\delta\mathbf{r} = 0$, $J = 1$, and a fixed M selecting a known two-orbital bare Bloch Slater determinant. The four local-energy components are then compared with an independent finite-basis evaluation. This test uses 4,096 walkers, 120 burn sweeps and 10 measurements and is never used for an optimization update.

Table 4: Fixed-state Hamiltonian regression in meV for the complete many-body state. Parentheses are Monte Carlo standard errors.

term	finite-basis	Monte Carlo	difference
kinetic	39.4207	39.3467(3343)	-0.0740
moiré	-38.5168	-38.7739(5347)	-0.2571
tunnelling	-29.9435	-29.6118(2984)	+0.3317
Coulomb	-9.6769	-10.1671(3354)	-0.4902
total	-38.7164	-39.2060(3354)	-0.4896

Every component and the total agree within 1.5 quoted standard errors. The regression therefore finds no common Hamiltonian normalization, Coulomb background, tunnelling, or energy-zero mismatch at its present precision.

A targeted regression test in `tests/test_jax_neural_bloch.py` instruments the real `_encoded_graph` function. One local-energy call produces four complete graph evaluations: base X with shape $(B, N_e, 2)$, all plus and all minus configurations with shape $(2N_e B, N_e, 2)$ each, and all one-particle layer flips with shape $(N_e B, N_e, 2)$. It checks every displaced coordinate and every flipped layer against an independently constructed array. The test passes, as does the JAX-PyTorch full-local-energy parity test to 2×10^{-7} meV absolute tolerance. No Hamiltonian input-routing error was found in this audit.

15 MCMC and natural-gradient optimization

There are 4,128 simultaneous walkers. Initial fractional coordinates are uniform on the torus and layer indices are uniform in $\{0, 1\}$. Each sweep proposes, for every walker, one random-particle Gaussian displacement with standard deviation 0.06 in fractional supercell coordinates, followed by one random-particle layer flip. Both use the Metropolis ratio $\min[1, |\Psi(X')/\Psi(X)|^2]$. The present jobs use 300

burn-in sweeps and two sweeps per update. Acceptance for coordinate and layer moves is recorded separately.

For real parameters θ_a ,

$$O_a(X) = \partial_{\theta_a} \log \Psi_{\theta}(X), \quad (58)$$

$$S_{ab} = \text{Re} [\langle O_a^* O_b \rangle - \langle O_a^* \rangle \langle O_b \rangle], \quad (59)$$

$$F_a = 2 \text{Re} \langle (E_{\text{loc}} - \langle E_{\text{loc}} \rangle) O_a^* \rangle. \quad (60)$$

The update is

$$(S + \lambda I) \Delta\theta = F, \quad \theta \leftarrow \theta - \eta \Delta\theta, \quad (61)$$

with ordinary, unpreconditioned CG. JAX JVP/VJP operations apply the metric matrix-free. Starting at iteration 20 and after every subsequent Krylov step, the code explicitly recomputes the authoritative residual

$$r_{\text{rel}}^{\text{true}} = \frac{\|F - (S + \lambda I) \Delta\theta\|_2}{\|F\|_2} \quad (62)$$

with an additional Fisher-vector product. The solve stops immediately when $r_{\text{rel}}^{\text{true}} < 0.02$ or at 60 CG iterations. At the cap, an update is accepted only if $r_{\text{rel}}^{\text{true}} < 0.05$ and $F^T \Delta\theta > 0$. Otherwise the parameter update is rejected and the following step uses the recovery values $\lambda = 0.03$ and $\eta = 10^{-3}$. After two consecutive stable recovery solves, the protocol returns to $\lambda = 0.01$ and $\eta = 2 \times 10^{-3}$. Every trace stores the true residual, $F^T \Delta\theta$, force and direction norms, acceptances, and whether the update was applied. This is an adaptive truncated natural gradient, not a claim that the 10^{-3} linear-system tolerance has been reached.

Table 5: Optimization and numerical settings. Quantities not specified by Ref. [1] are explicitly marked as local choices.

quantity	current value	provenance
width / message iterations	32 / 2, unshared	Ref. [1], Appendix B
orbital MLP hidden width	32, GELU	Ref. [1], Appendix B
final determinants	1	local controlled choice
optimizer / implementation	natural gradient / JAX 0.11.1	Ref. [1]
learning rate	2×10^{-3}	Ref. [1]
walkers per update	4,128	Ref. [1]
affine-kernel initialization	LeCunNormal	Ref. [1]
h^0 initialization	random, trainable	random: Ref. [1]; law/sharing: local
full- M initialization	complex LeCunNormal	user choice; paper unspecified
node aggregation	$\sum_j m_{ij}$, including self	Ref. [1], Eq. (19)
updates in present experiment	120	user-specified finite run
paper long-run statement	1,000 for clusters $> 3 \times 3$	Ref. [1]
burn / sweeps per update	300 / 2	local choice
proposal scale	0.06 fractional	local choice
SR damping	10^{-2}	local choice
CG range / true-residual rules	20-60 / 0.02 stop, 0.05 accept	user-specified
wavefunction / SR batch	516 / 516	local memory choice
local-energy batch	32	local memory choice
finite-difference step	$1.5 \times 10^{-3} a_M$	local choice
random seed	83	matched local choice
precision	JAX 64 bit	local numerical choice

All 4,128 walkers are used in the energy, force, and metric estimate at each update. The printed within-step SEM treats walker values as an instantaneous ensemble diagnostic; it is not an independent uncertainty. Checkpoints at initialization, every 10 updates, the latest completed update, and update 120 store both flattened parameters, optimizer recovery state, sampler RNG, and walker coordinates/layers. Interrupted jobs resume this complete state; no validation configuration ever enters an update.

16 Implementation checks completed before production

The current complete 51-test suite passes; selected architecture, parity, observable, and local-energy checks verify:

1. particle exchange changes the determinant sign and no other factor;
2. J_{ki} is applied exactly once;
3. fixed Bloch coefficients are absent from trainable parameters;
4. shifted-coordinate Bloch evaluation agrees with an explicit plane-wave calculation;
5. PyTorch and JAX $\log \Psi$ agree at identical parameters;
6. complete JAX and PyTorch local energies agree;
7. no- M is invariant under both primitive translations and contains no trainable M array;
8. the complete-state C_3 projector satisfies character, idempotency, orthogonality, nonzero-sector-norm and projector-disabled no- M regression tests before its three-sector production run;
9. matrix-free force and Fisher-vector products agree with an explicit score-Jacobian calculation;
10. checkpoint export between JAX and PyTorch is lossless;
11. full- M is reproducibly random, complex, and not one-hot;
12. h_i^0, h_{ij}^0 are nonzero trainable parameters; and
13. a one-particle graph receives its nonzero $j = i$ message in Eq. (19);
14. the 1-RDM uses full replaced-configuration ratios and a fixed bare projector, and its bare-Slater full-basis trace is N_e ; and
15. kinetic displacements and tunnelling flips rebuild the complete JAX graph, while diagonal Hamiltonian factors use raw coordinates.

A one-update end-to-end smoke test also passed MCMC, local-energy evaluation, truncated SR, residual recording, parameter update, and checkpoint writing. These checks validate implementation consistency, not physical convergence.

17 Controlled v3 sector-mixer experiments

The historical full- M jobs provide the unrestricted baseline. The two v3 no- M jobs train M_S for 120 updates with 4,128 persistent walkers. For each filling, outer- C_3 first performs 20 shared raw-state updates, copies only the learned parameters into three fresh chains, and then performs 100 projected updates independently for $m = 0, 1, 2$. Walker and optimizer states are not shared between branches. Public sector ordering uses the final-ten training mean.

Table 6: Authoritative v3 calculations used in this report. Outer branches use 20 shared raw updates followed by 100 independently sampled projected updates.

ν	ansatz	K_{tot}	result-directory suffix
1/3	full M	mixed	adaptive_cg120_20260821_full_m_nu1of3
1/3	no $M+M_S$	Γ	sector_mixer_v3/no_m/nu1of3
1/3	outer P_0+M_S	$\Gamma, m = 0$	sector_mixer_v3/outer_c3/nu1of3/m0
1/3	outer P_1+M_S	$\Gamma, m = 1$	sector_mixer_v3/outer_c3/nu1of3/m1
1/3	outer P_2+M_S	$\Gamma, m = 2$	sector_mixer_v3/outer_c3/nu1of3/m2
2/3	outer P_0+M_S	$\Gamma, m = 0$	sector_mixer_v3/outer_c3/nu2of3/m0
2/3	outer P_1+M_S	$\Gamma, m = 1$	sector_mixer_v3/outer_c3/nu2of3/m1
2/3	outer P_2+M_S	$\Gamma, m = 2$	sector_mixer_v3/outer_c3/nu2of3/m2
2/3	full M	mixed	adaptive_cg120_20260821_full_m_nu2of3
2/3	no $M+M_S$	Γ	sector_mixer_v3/no_m/nu2of3

17.1 Training-chain diagnostics

Table 7 summarizes the final ten attempted updates for each frozen trace. The mean and RMS describe the persistent training chain and are not independent energy estimates. A rejected update contributes its sampled local energy to the diagnostic trace but leaves the parameters unchanged. All three v3 $\nu = 1/3$ outer branches accepted every update. At $\nu = 2/3$, the v3 outer P_0, P_1, P_2 branches accumulated 0, 5 and 0 rejected updates, respectively.

Table 7: Final-ten training diagnostics. Energies are meV per particle; RMS is training-chain scatter.

ν	ansatz	$\langle E/N_e \rangle$	min. E/N_e	RMS	$\langle r_{\text{rel}} \rangle$	$\langle N_{\text{CG}} \rangle$	reject ₁₂₀
1/3	full M	-35.9952	-36.0866	0.0434	0.0183	48.0	0
1/3	no $M+M_S, \Gamma$	-37.3902	-37.4468	0.0287	0.0176	31.9	0
1/3	outer P_0+M_S	-37.3458	-37.4276	0.0516	0.0155	20.8	0
1/3	outer P_1+M_S	-32.4638	-32.5541	0.0570	0.0169	43.2	0
1/3	outer P_2+M_S	-36.0804	-36.1597	0.0465	0.0147	26.9	0
2/3	outer P_0+M_S	-52.5241	-52.5965	0.0620	0.0179	39.8	0
2/3	outer P_1+M_S	-51.5901	-51.6787	0.0678	0.0179	48.2	5
2/3	outer P_2+M_S	-52.8211	-52.9017	0.0477	0.0171	38.7	0
2/3	full M	-52.0575	-52.1563	0.0817	0.0371	54.5	19
2/3	no $M+M_S, \Gamma$	-52.7841	-52.8535	0.0394	0.0253	58.5	4

The v3 no- M final-ten position/layer acceptances are 0.674/0.637 and 0.656/0.611 at $\nu = 1/3$ and $2/3$. At $\nu = 1/3$, outer P_0, P_1, P_2 give 0.676/0.628, 0.696/0.594 and 0.678/0.537, with mean local-energy variances 53.2, 114.3 and 72.4 meV². At $\nu = 2/3$, the P_0, P_1, P_2 tails give 0.653/0.608, 0.660/0.621 and 0.650/0.613, with variances 690.5, 725.7 and 667.6 meV². Thus a smooth energy trace does not establish a zero-variance eigenstate.

17.2 Training-tail comparison

Following the optimization convention of Ref. [1], the public energy comparison now uses the equilibrated persistent training ensemble. For each trace we average the final ten measurements and quote their RMS fluctuation; the RMS is a correlated training-chain stability diagnostic, not an independent standard error. Among the three outer sectors, P_0 is lowest at $\nu = 1/3$, with $\langle E/N_e \rangle_{\text{tail}} = -37.3458$ meV and $\Delta E_{5b} = +0.0474$ meV. At $\nu = 2/3$, P_2 is lowest, with $\langle E/N_e \rangle_{\text{tail}} = -52.8211$ meV and $\Delta E_{5b} = -0.0956$ meV. The latter is an optimization-chain result, not yet a multi-seed variational bound.

Figure 7 retains all three m sectors outside Fig. 6. Its upper row plots their final-ten energy offsets from the lowest sector; the lower row reports rejected parameter updates directly, replacing the black failure crosses formerly overlaid on the CG traces. At $\nu = 1/3$, all displayed v3 runs have zero rejected updates. A rejected update leaves the parameters unchanged but its sampled local energy remains part of the diagnostic trace.

18 Post-training one- and two-body observables

Figure 8 compares full- M and v3 no- $M+M_S$ at both fillings. Figures 9 and 10 display all three externally projected sectors. At $\nu = 1/3$, the first row is the training-energy minimum P_0 and is the selected out- C_3 result; P_1 and P_2 are retained only for a controlled sector comparison. At $\nu = 2/3$, P_2 is the corresponding minimum. The panels are regenerated from the equilibrated final training ensembles. Each contains 4,128 configurations; energy selection and observable estimation use the same production-chain convention as the optimization protocol of Ref. [1].

18.1 Full-ratio 1-RDM and fixed bare-band projector

For $a = (n, \mathbf{k})$ and $b = (m, \mathbf{q})$, the estimator is

$$\begin{aligned} \rho_{ab} &= \langle c_b^\dagger c_a \rangle \\ &= 2N_e \mathbb{E}_{X \sim |\Psi|^2, i, x'} \left[(\phi_a^0(x_i))^* \phi_b^0(x') \left(\frac{\Psi(X_{i \rightarrow x'})}{\Psi(X)} \right)^* \right]. \end{aligned} \quad (63)$$

The factor two accounts for uniform auxiliary sampling of the two layers. In every ratio, replacing one particle triggers a complete new ansatz call: the code rebuilds the graph and recomputes all particles' $\delta \mathbf{r}_j$, $J_{\mathbf{k}j}$, and either the MQ determinant or all M_S -weighted sector minors. The projector is instead the fixed bare function

$$\phi_{n\mathbf{k}}^0(\mathbf{r}, l) = \sum_{\mathbf{G}} u_{n\mathbf{k}\mathbf{G}}^l e^{i(\mathbf{k}+\mathbf{G}) \cdot \mathbf{r}}; \quad (64)$$

no backflow phase is multiplied into it. Thus neural dressing appears once, only through the complete wavefunction

ratio.

All 4,128 equilibrated training configurations and 16 auxiliary points per configuration are used to estimate the 45×45 five-band-projected matrix. The band populations are

$$W_n = \frac{1}{N_e} \sum_{\mathbf{k}} \rho_{n\mathbf{k}, n\mathbf{k}}. \quad (65)$$

Table 8: Five-band-projected 1-RDM results from equilibrated training samples. Band weights are percentages; parentheses are eight-block estimates. ED weights use the exact five-band one-body density matrix.

ν	state	trace	W_1	W_2	W_3	W_4	W_5	min eig.
1/3	five-band ED	3	96.47	2.95	0.52	0.05	0.01	—
1/3	no $M+M_S$	2.890(98)	95.53(86)	2.75(1.27)	-1.57(89)	0.16(56)	-0.53(82)	-0.084
1/3	outer P_0+M_S	3.077(105)	97.18(96)	4.70(87)	0.25(1.13)	-0.37(1.17)	0.79(1.41)	-0.080
1/3	outer P_1+M_S	3.019(67)	83.34(95)	14.99(42)	1.57(89)	0.96(91)	-0.23(73)	-0.098
1/3	outer P_2+M_S	2.917(95)	96.73(1.13)	1.76(1.52)	-1.36(1.95)	0.76(58)	-0.68(88)	-0.142
2/3	five-band ED	6	88.95	7.59	2.53	0.76	0.16	—
2/3	no $M+M_S$	5.894(113)	89.03(74)	6.19(1.11)	1.21(89)	0.80(1.08)	1.01(95)	-0.196
2/3	outer P_0+M_S	6.280(109)	91.93(74)	7.28(1.02)	3.05(75)	1.79(80)	0.61(1.35)	-0.169
2/3	outer P_1+M_S	5.879(140)	86.97(95)	7.98(79)	1.94(1.49)	0.83(1.25)	0.27(72)	-0.318
2/3	outer P_2+M_S	5.853(214)	90.19(1.63)	5.13(1.57)	1.21(1.10)	-0.32(77)	1.33(1.14)	-0.354

Most projected traces are compatible with N_e at the stated precision, but the matrices are only Hermitized, not projected onto the positive-semidefinite cone. Their negative minimum eigenvalues therefore forbid a stable natural-orbital eigenspace interpretation. The diagonal band populations are more robust and reveal a real sector dependence. At $\nu = 1/3$, P_0 and P_2 retain about 97% band-1 weight, whereas P_1 has 83.34(95)% and transfers 14.99(42)% to band 2; five-band ED has $W_1 = 96.47\%$ and $W_2 = 2.95\%$. At $\nu = 2/3$, all three sectors are closer to the ED distribution (88.95, 7.59, 2.53, 0.76, 0.16)%; the lowest-training P_2 gives (90.19, 5.13, 1.21, -0.32, 1.33)%. Small negative entries remain consistent with auxiliary-coordinate noise.

The estimator now caches $\Psi_m(X)$ once per configuration block and jointly batches all 16 replacement coordinates. Each replacement still rebuilds the complete graph, all $\delta \mathbf{r}_j$, J , the M_S determinant mixture, the projected three-branch sum and the wavefunction ratio. A 128-configuration, four-draw probe takes 3.0 seconds; running two 32-thread jobs concurrently completes all six 4,128-by-16 sector measurements in three batches instead of repeating the former nested small-call path. The focused bare-Slater regression sets $\delta \mathbf{r} = 0$, $J = 1$, fixes a complex 2×9 matrix M with orthonormal rows, and samples the resulting single determinant. In a complete nine-orbital test basis the exact matrix is $\rho = M^T M^*$, with trace two. With 4,096 walkers, 100 burn sweeps and 32 auxiliary draws, the measured trace is 2.01047 ± 0.01459 , the Frobenius error is 0.05272, the Hermiticity residual is zero, and the minimum sampled eigenvalue is -0.02072. This unit test verifies occupation, Hermiticity, PSD within sampling noise and full-basis trace normalization.




Figure 6: Controlled neural-Bloch optimization, multiband energy, band mixing and CG stability. Top row: $\nu = 1/3$; bottom row: $\nu = 2/3$. Blue and red denote full M and v3 no $M+M_S$; orange denotes the lowest final-ten v3 outer sector, P_0 at $\nu = 1/3$ and P_2 at $\nu = 2/3$. (a,e) Persistent-walker training energies. (b,f) One- through five-band ED energies and final-ten training means; shaded widths are the corresponding RMS, not independent error bars. (c,g) Bare-band populations from the complete-wavefunction-ratio 1-RDM, including the lowest-training outer sector and five-band ED in gray. (d,h) Explicit true CG residuals; horizontal lines mark the 0.02 early-stop and 0.05 cap-acceptance thresholds. Rejected-update counts are summarized separately in the text.

18.2 Full structure factor, connected part, and density order

The structure factor is evaluated directly as a two-body observable, not reconstructed from the 1-RDM:

$$S_{\text{full}}(\mathbf{q}) = \frac{\langle |\rho_{\mathbf{q}}|^2 \rangle}{N_e}, \quad (66)$$

$$S_c(\mathbf{q}) = \frac{\langle |\rho_{\mathbf{q}}|^2 \rangle - |\langle \rho_{\mathbf{q}} \rangle|^2}{N_e}, \quad (67)$$

$$m_{\mathbf{q}} = \frac{|\langle \rho_{\mathbf{q}} \rangle|}{N_e}. \quad (68)$$

The ED comparison uses S_{full} with the same normalization. At $\nu = 1/3$, outer P_0, P_1, P_2 give maximum nonzero-momentum $S_{\text{full}} = 1.4845(131), 1.1205(120), 1.2687(126)$, respectively, compared with the five-band ED value 1.4381. At $\nu = 2/3$ the corresponding values are 0.7585(51), 0.6944(145), 0.7059(101), compared with ED 0.6197. The sector dependence is therefore visible in both the 1-RDM and the genuine two-body observable; it is not a plotting artifact.

For $m = 0, 1, 2$, the largest nonzero- q density amplitudes are 0.00385(418), 0.00895(399), 0.01527(465) at

$\nu = 1/3$ and 0.00322(242), 0.00240(361), 0.00580(197) at $\nu = 2/3$. These small values show negligible static translation-breaking expectation at the sampled reciprocal vectors; the large S_{full} peaks are density fluctuations. The real-space panels in Figs. 9 and 10 include the intra-moiré-cell Bloch density and should not be identified with a nonzero CDW order parameter by visual contrast alone.

19 Conclusion

The implemented baseline follows the printed Luo–Fu Appendix-B structure: fixed mean-field band-1 plane-wave coefficients, periodic node and edge features, trainable random auxiliary features, two unshared width-32 message iterations with the literal self-inclusive $\sum_j m_{ij}$, coordinate backflow, direct complex $J_{\mathbf{k}i}$ applied exactly once, and one final generalized Slater construction. Full M mixes momenta at the single-particle level and therefore generally mixes total-momentum sectors. Version 3 no- M instead uses $M_S \in \mathbb{C}^{1 \times 12}$ to combine all and only the selected-momentum determinants in the exact total- Γ sector.




Figure 7: Outer- C_3 sector splitting and optimizer stability. (a,b) Final-ten energy offsets of the externally projected P_m states, measured from the lowest sector at each filling. Error bars are RMS fluctuations of the correlated training tail, not independent standard errors. At $\nu = 1/3$, $m = 0$ is selected, while $m = 1, 2$ lie 4.882 and 1.265 meV per particle higher. At $\nu = 2/3$, $m = 2$ is selected, with $m = 0, 1$ higher by 0.297 and 1.231 meV per particle. The smaller spread at $\nu = 2/3$ is compatible with, but does not establish, the expected quasi-degenerate FCI manifold. (c,d) Cumulative rejected parameter updates for the methods retained in Fig. 6; rejection concerns an optimizer step, not a Monte Carlo configuration.

The outer construction leaves that raw network unrestricted and projects the complete M_S determinant mixture with P_m . Character and projector-algebra residuals are below 6.4×10^{-14} and 7.6×10^{-15} , all three projected norms are nonzero, and local energy, Metropolis ratios and SR derivatives all use the projected state. It adds no trainable parameter to the 78,124-parameter v3 no- M model.

On the final-ten training tail, v3 no- M reaches -37.3902 meV per particle at $\nu = 1/3$, while outer P_0 is the lowest of the three projected sectors at -37.3458 meV. At $\nu = 2/3$, v3 no- M reaches -52.7841 meV and outer P_2 is the lowest training sector at -52.8211 meV, 0.0956 meV below the five-band ED benchmark. This promising crossing is a correlated, single-seed optimization-chain result and is not promoted to an independent variational bound.

The all-sector observables show that $m = 0, 1, 2$ cannot be represented by one duplicated map: at $\nu = 1/3$, P_1 has markedly enhanced band-2 weight and a smaller structure-factor peak, while at $\nu = 2/3$ the differences are weaker but remain resolved. Thus the complete-state projector changes both energy and correlations, rather than merely relabeling the same generalized determinant. The complete-ratio 1-RDM and direct two-body structure factors remain useful diagnostics, but their non-PSD noise and the single-seed finite optimization preclude a natural-orbital or topological claim. The next controlled tests are longer multi-seed optimization and sector-selection replication.

References

- [1] D. Luo, T. Zaklama, and L. Fu, “Solving fractional electron states in twisted MoTe_2 with deep neural network,” arXiv:2503.13585v3, doi:10.48550/arXiv.2503.13585.
- [2] X. Li, Y. Chen, B. Li, H. Chen, F. Wu, J. Chen, and W. Ren, “Deep Learning Sheds Light on Integer and Fractional Topological Insulators,” arXiv:2503.11756v1, doi:10.48550/arXiv.2503.11756.
- [3] G. Pescia, J. Nys, J. Kim, A. Lovato, and G. Carleo, *Message-Passing Neural Quantum States for the Homogeneous Electron Gas*, arXiv:2305.07240v3, doi:10.48550/arXiv.2305.07240.
- [4] D. Luo and B. K. Clark, “Backflow transformations via neural networks for quantum many-body wave functions,” Phys. Rev. Lett. **122**, 226401 (2019), doi:10.1103/PhysRevLett.122.226401.
- [5] A. P. Reddy *et al.*, “Fractional quantum anomalous Hall states in twisted bilayer MoTe_2 and WSe_2 ,” Phys. Rev. B **108**, 085117 (2023), doi:10.1103/PhysRevB.108.085117.
- [6] Z. Lin *et al.*, “Fractional Chern insulator states in an isolated flat band of zero Chern number,” Newton **2**, 100339 (2026), doi:10.1016/j.newton.2025.100339.
- [7] L. Zhang, G. Duan, and D. Luo, “WF-Bench: A benchmark for neural network wavefunction expres-

sivity and scaling laws," arXiv:2605.29683 (2026),
[doi:10.48550/arXiv.2605.29683](#).

[8] M. Yamazaki *et al.*, "Physics-inspired transformer quantum states via latent imaginary-time evolution," Phys. Rev. Research **8**, 033032 (2026),
[doi:10.1103/bjxb-8tsk](#).

[9] R. Rende, L. Viteritti, and A. Georges, "Superconductivity in the t - t' Hubbard model from symmetry-preserving neural-network quantum states," arXiv:2608.12465 (2026),
[doi:10.48550/arXiv.2608.12465](#).




Figure 8: Full- M and $v3$ no- $M+M_S$ observables. The first two rows are $\nu = 1/3$ and the last two are $\nu = 2/3$; at each filling full M is followed by no $M+M_S$. Columns show bare-band-1 momentum occupation, the full static structure factor, and folded charge density. The $v3$ no- M panels use the newly evaluated 4,128-configuration final ensembles.




Figure 9: $\nu = 1/3$ outer-sector observables from equilibrated training samples. Rows are complete-wavefunction projectors P_0, P_1, P_2 ; the first row, P_0 , is explicitly marked as the selected outer- C_3 state because it has the lowest final-ten training energy. Columns show bare-band-1 momentum occupation from the complete-ratio 1-RDM, the full static structure factor with ED normalization on the complete hexagonal BZ, and folded charge density over a two-by-two moiré-cell parallelogram. Shared color limits expose the strong P_1 band-mixing and structure-factor difference without rescaling individual sectors.




Figure 10: $\nu = 2/3$ outer-sector observables from equilibrated training samples. Rows are P_0, P_1, P_2 and columns use the same definitions and global color limits as Fig. 9. All panels use 4,128 configurations and 16 jointly batched auxiliary replacements per configuration. The three sectors are shown separately because both their band populations and two-body structure factors differ beyond a simple identical-map interpretation. The P_2 row is marked as the lowest final-ten-energy sector.